

EmbedGuard: Cross-Layer Detection and Provenance Attestation for Adversarial Embedding Attacks in RAG Systems

Author: Neeraj Kumar Singh Beshane

Independent Researcher, California, USA

Corresponding Author: b.neerajkumarsingh@gmail.com

Competing Interests: The author is an employee of Parafin. This work was conducted independently and is not affiliated with the author's employer.

Post-publication manuscript version: 3.2

Ethics Statement: This research evaluates defensive AI-security mechanisms using synthetic attack strings and public benchmark queries. It involves no human participants, private personal data, or testing against live third-party systems. No institutional review was sought. Released attack artifacts are limited to inputs needed to reproduce the detector benchmark and are paired with the corresponding defensive implementation.

Abstract

Embedding-based Retrieval-Augmented Generation (RAG) systems remain vulnerable to corpus poisoning and prompt injection across multiple architectural stages. EmbedGuard proposes a four-layer reference architecture combining prompt detection, embedding provenance, retrieval-distribution monitoring, output-consistency analysis, and a deterministic correlation decision. The target provenance protocol uses an AMD SEV-SNP trust root, while the released package provides only an HMAC software simulation of that protocol; retrieval and output modules are experimental prototypes. The open Tier-2 regression benchmark (N=135) exercises only the 83-pattern prompt detector and observes 30/30 included attacks detected and 0/105 benign queries flagged, with two-sided 95% Wilson intervals of 88.6%-100% for detection and 96.5%-100% for specificity. These fixed, mostly one-example attack categories measure regression-set coverage rather than generalization, hardware attestation, or cross-layer benefit. A second open benchmark (Tier 2b) injects the released PoisonedRAG adversarial passages into a BEIR Natural Questions retrieval pool. Provenance attestation accepted every poisoned passage submitted through the approved ingestion path (0/50 detected) because it certifies origin, not content. The retrieval layer, restricted to its distribution-distance component and calibrated on clean traffic only, flagged 50/50 black-box and 48/50 stealth attacks with five planted passages per target, at 6/200 flagged clean queries; detection fell to 43/50 and 33/50 with three passages and to 7/50 with one. With the package's default settings it detected none. The IJCESEN version of record separately reports a production-scale four-layer evaluation and an 18.4 percentage-point ablation gain, but the raw predictions, production corpus, baseline configurations, and hardware-attestation logs for that evaluation are not contained in this repository; this revision therefore treats those numbers as archived publication claims rather than independently reproduced evidence. The implementation, Tier-2 inputs, Tier-2b fetch and run scripts, and per-query results are released under Zenodo concept DOI 10.5281/zenodo.18364919.

Keywords: Retrieval-Augmented Generation Security, Corpus Poisoning, Embedding Space Poisoning, Cross-Layer Attack Detection, Trusted Execution Environments, Provenance Attestation

1. Introduction

With the advent of large language models and their deployment in enterprise applications, Retrieval-Augmented Generation (RAG) systems have emerged as one of the most impactful architectures for artificial intelligence applications. RAG systems combine the generative capabilities of neural language models with the ability to retrieve information dynamically from external knowledge sources, alleviating critical drawbacks of purely generative models such as knowledge staleness, factual hallucinations, and limited domain coverage (Lewis et al., 2020). This architectural pattern has become ubiquitous in production deployments across healthcare, financial services, legal research, and customer service applications.

Recent security research has identified critical vulnerabilities in RAG retrieval components, particularly embedding space poisoning attacks where adversaries insert maliciously constructed documents into the retrieval knowledge base to influence the generation process (Zou et al., 2025; Liu et al., 2023). These attacks exploit high-dimensional embedding geometry: even minimal corpus contamination (less than 1% of documents) can achieve attack success rates exceeding 80% through strategic semantic space positioning. Research demonstrates that attackers can generate documents that meet retrieval targets for specific query patterns while remaining sufficiently semantically diverse to evade clustering-based outlier detection techniques (Zou et al., 2025). The permanence of embedding attacks differentiates them from transient prompt-based exploits, combining supply chain attack stealth with runtime exploit immediacy to create a distinct and persistent threat surface.

1.1 A Live Production Threat, Not a Hypothetical

This threat class has moved from academic demonstration to documented production impact. CVE-2025-32711 (EchoLeak) against Microsoft 365 Copilot demonstrated a network-reachable, zero-click AI command injection with high confidentiality impact in a flagship enterprise assistant (NVD, 2025). ConfusedPilot showed that malicious documents placed inside an enterprise network could alter Microsoft 365 Copilot's responses, suppress legitimate documents, and violate sharing boundaries (RoyChowdhury et al., 2024). Phantom demonstrated that a single poisoned document could compromise a commercial RAG application (NVIDIA ChatRTX) black-box, achieving 48% success for passage exfiltration (Chaudhari et al., 2024). Persistent-memory variants of the same pattern have been demonstrated against ChatGPT and Gemini, where injected content in retrievable memory drives continuous data exfiltration (Rehberger, 2024; Rehberger, 2025).

Industry threat frameworks now name this class formally. OWASP's Top 10 for LLM Applications 2025 dedicates LLM08 to Vector and Embedding Weaknesses — exploitation of RAG embedding stores to "inject harmful content, manipulate model outputs, or access sensitive information" — alongside LLM01 (Prompt Injection, explicitly covering indirect injection via retrieved content) and LLM04 (Data and Model Poisoning, explicitly covering manipulated embedding data) (OWASP, 2025). MITRE ATLAS catalogues the attack kill chain with dedicated techniques: AML.T0064 (Gather RAG-Indexed Targets), AML.T0066 (Retrieval Content Crafting), AML.T0070 (RAG Poisoning), and AML.T0071 (False RAG Entry Injection) (MITRE, 2025). NIST's adversarial machine learning taxonomy (AI 100-2e2025) classifies RAG knowledge-base poisoning under NISTAML.013 and indirect prompt injection under NISTAML.015 (Vassilev et al., 2025).

The exposed surface is broad: Menlo Ventures' 2024 survey of 600 U.S. enterprise leaders found RAG at 51% adoption among enterprise generative-AI deployments, up from 31% the prior year (Menlo Ventures, 2024). Half of enterprise AI deployments inherit an attack surface for which named adversary techniques exist, real CVEs have shipped, and single-stage defenses have documented evasions (Section 2.5).

1.2 Economic and Security Implications

These vulnerabilities have substantial economic implications for organizations deploying RAG systems. IBM Security's 2024 Cost of a Data Breach Report places the global average cost of a breach at USD 4.88 million and the mean time to identify and contain one at 258 days (IBM Security, 2024). The report does not isolate breaches of AI or RAG systems, and it finds that organizations making extensive use of security AI and automation had lower average costs (USD 3.84 million, against USD 5.72 million for organizations not using them), so these figures describe breaches in general rather than embedding-space attacks. The persistence of embedding-space attacks exacerbates these costs, as poisoned vectors remain in knowledge bases until manually identified and removed, resulting in prolonged compromise timeframes. This permanence, combined with the difficulty of forensic analysis in high-dimensional embedding spaces, creates extended uncertainty regarding breach scope and impact.

The high-dimensionality of embedding spaces (typically 768 to 1536 dimensions for modern embedding models) enables adversaries to construct documents that preserve semantic relevance for target query patterns while remaining grammatically valid and linguistically coherent, thus evading perplexity-based statistical detectors. Furthermore, adversarial embeddings demonstrate transferability between embedding models, meaning attackers who optimize attacks against publicly available models can successfully transfer them to proprietary models with high confidence of success (Zou et al., 2025; Xiang et al., 2024).

1.3 Limitations of Current Defense Mechanisms

Contemporary defense mechanisms primarily adopt stage-specific approaches, optimizing detection for isolated attack surfaces within the RAG architecture. RAGuard employs a two-layer defense combining adversarial retriever training with chunk-wise perplexity filtering and text similarity analysis (Cheng et al., 2025). RobustRAG implements isolate-then-aggregate strategies with certifiable guarantees using keyword-based voting (Xiang et al., 2024). TrustRAG uses K-means cluster filtering with LLM self-assessment for malicious document detection (Zhou et al., 2025a). These defenses may use several stages, but they do not combine anomaly evidence from the prompt, provenance, retrieval, and output stages into one decision. Whether such a combination improves detection against attacks spread across stages is an open empirical question that this work does not settle.

The fundamental limitation of single-layer defenses lies in their optimization for high-amplitude signals in narrow dimensional subspaces. Perplexity-based filters assume poisoned documents exhibit linguistic incoherence, yet advanced adversaries generate fluent malicious text indistinguishable from legitimate documents. Clustering-based methods assume poisoned embeddings appear spatially anomalous, yet attackers optimize for embedding centrality while maintaining target query similarity. Activation-based methods assume poisoned content causes abnormal model behavior, yet adversaries craft documents producing contextually appropriate activation patterns. Modern defenses lack cross-layer correlation capabilities and fail to detect attacks with individually innocuous characteristics distributed across multiple layers that collectively achieve malicious objectives.

1.4 Contributions

To address these limitations, we present EmbedGuard, a cross-layer reference architecture with a target hardware-attestation design for RAG systems. Layered RAG defenses have begun to appear — frameworks that stack prompt-injection filters (Saleem et al., 2026), orchestrate independent per-layer defenses (Pallerla et al., 2026), or combine embedding anomaly detection with response verification (Ramakrishnan & Balaji, 2025) — and software-signature schemes for embedding provenance have been proposed (Wanger, 2026). EmbedGuard differs from these on two axes: it proposes fusion of anomaly signals from all four RAG architectural layers into one decision rather than independent filters, and it proposes moving embedding-provenance signing into a hardware root of trust. The released package implements the fusion logic and a software HMAC simulation of the provenance protocol; it does not implement AMD SEV-SNP attestation. The framework makes the following contributions:

Cross-Layer Signal Fusion: The literature scan in Section 2.4 found close layered systems but no prior system with this exact four-signal fusion and TEE-rooted provenance combination. That is a scoped integration claim, not a claim that layered RAG defense or provenance was invented here. The current deterministic engine combines confidence-weighted consensus, a strongest-signal floor, and a multi-layer correlation boost.

Hardware-Rooted Provenance Design: We specify a target AMD SEV-SNP protocol that binds document, model, vector, and platform measurements. If implemented with real attestation verification, this would provide integrity and provenance against post-ingestion vector tampering or unauthorized embedding insertion. It would not prove that an authorized source document is benign or true; malicious inputs admitted through an approved ingestion path can still receive valid provenance.

Reproducible Benchmark Evaluation: Evaluation on the EmbedGuard open regression set—comprising locally curated Natural Questions-style, HotpotQA-style, and MS-MARCO-style benign files plus a curated 25-category injection file—observes 30/30 included attacks detected and 0/105 included benign queries flagged for the released pattern-only prompt detector. The 95% Wilson intervals are 88.6%-100% for detection and 96.5%-100% for specificity. The named benign files lack upstream IDs and extraction manifests, and mostly one-sample attack categories do not establish generalization. Complete code, inputs, and evaluation scripts are released with one-command repeatability:

```
git clone https://github.com/neerazz/embedguard && cd embedguard && ./reproduce.sh
```

Open Corpus-Poisoning Evaluation: Section 4.3 evaluates the provenance and retrieval layers against the published PoisonedRAG passages on BEIR Natural Questions, with thresholds calibrated only on clean traffic, a held-out test split, and a control that runs the target questions against the unpoisoned pool. It shows that provenance attestation does not detect poisoned passages admitted through the approved ingestion path, and that a calibrated distribution-distance statistic detects poisoning that fills most of the retrieved set but not a single planted passage.

Attack Coverage Analysis: The fixed regression set contains 25 attack categories including direct injection, jailbreak attempts, encoding obfuscation (Unicode, Base64), delimiter confusion, XML/Markdown injection, hypothetical/fictional framing, translation attacks, authority claims, emotional manipulation, RAG-specific attacks, composite multi-vector attacks, and subtle manipulation. All included samples are detected with scores from 0.80 to 1.0; mostly one-sample categories do not establish category-level generalization.

Decision Policy Modes: Three modes map scores to LOG, FLAG, or BLOCK decisions. The package returns those decisions; application code must implement logging, review queues, blocking, and fallbacks.

The remainder of this paper is organized as follows: Section 2 provides background on RAG security and related work. Section 3 details the EmbedGuard architecture and detection mechanisms. Section 4 presents experimental evaluation and comparative analysis. Section 5 discusses applications, limitations, and societal implications. Section 6 concludes with future research directions.

2. Background and Related Work

2.1 RAG Attack Surface and Poisoning Mechanics

The attack surface of RAG systems encompasses multiple architectural layers, each presenting distinct vulnerabilities that adversaries can exploit to manipulate system behavior. Knowledge poisoning attacks modify the retrieval mechanism, steering language models toward attacker-controlled content through careful manipulation of the embedding space and semantic similarity calculations fundamental to retrieval-based systems (Zou et al., 2025).

PoisonedRAG shows that five injected texts per target question, in corpora of millions of passages, are enough for high attack success (Zou et al., 2025). Adversaries generate documents that satisfy retrieval targets for specific query patterns while maintaining sufficient semantic diversity to evade clustering-based outlier detection. Its black-box variant prepends a copy of the target question to an LLM-generated passage, and its white-box variant refines that prefix with gradient-guided token substitution. The resulting passages are fluent, but they need not match the embedding distribution of benign documents; in the open evaluation of Section 4.3 they were separable by a distribution-distance statistic.

Table 1: RAG Attack Vectors and Poisoning Characteristics

Attack Component	Vulnerability Mechanism	Persistence Duration	Detection Complexity
Embedding Space Poisoning	Strategic document positioning in high-dimensional semantic space	Extended persistence until explicit removal	High complexity due to distributed vector storage
Gradient-Based Optimization	Iterative refinement maximizing retrieval probability	Sustained across query sessions	Difficult through traditional forensic techniques
Transferability Exploitation	Cross-architecture attack effectiveness	Long-term knowledge base compromise	Extended detection and containment timelines
Semantic Similarity Manipulation	Query-document matching exploitation	Persistent vector influence	Complex remediation requiring integrity validation

2.2 Economic Impact and Detection Challenges

Public breach-cost data do not yet separate incidents in AI retrieval systems from other breaches. IBM's 2024 report gives a mean of 258 days to identify and contain a breach and an average cost of USD 4.88 million across all breach types (IBM Security, 2024). We know of no published measurement of detection time for poisoned vector stores, so the argument that such poisoning is slow to detect rests on the persistence mechanism described in Section 2.1, not on breach statistics.

Cost analysis reveals that remediation expenses are highest when poisoning affects training data or model behavior, requiring poison purging, integrity validation, and potentially retraining in secure environments. Breaches affecting retrieval systems present additional recovery challenges due to distributed vector store architectures, where identifying all compromised embeddings at scale proves difficult. Forensic processes struggle to reason about attack impacts in high-dimensional embedding spaces, creating prolonged organizational uncertainty regarding breach scope.

2.3 Defense Mechanism Landscape

Contemporary defense mechanisms employ various strategies to protect RAG systems from poisoning attacks. RAGuard employs a two-layer approach combining adversarial retriever training with perplexity-based filtering and text similarity analysis at the retrieval layer (Cheng et al., 2025). RobustRAG implements isolate-then-aggregate strategies with certifiable guarantees, using keyword-based voting across retrieved documents (Xiang et al., 2024). TrustRAG uses K-means cluster filtering combined with LLM self-assessment for malicious document detection (Zhou et al., 2025a). More recently, ReliabilityRAG adopts a graph-theoretic perspective, identifying "consistent majority" documents through Maximum Independent Set computation on contradiction graphs with provable robustness guarantees (Shen et al., 2025). RAGDefender employs a post-retrieval approach using clustering-based grouping for single-hop queries and concentration-based analysis for multi-hop reasoning, achieving substantial attack success rate reductions (Kim et al., 2025a). PoisonedRAG established foundational attack methodologies demonstrating that even minimal corpus contamination achieves disproportionate attack success through strategic embedding space positioning (Zou et al., 2025). RevPRAG detects poisoned responses from the LLM's internal activations and reports a 98% true positive rate (Tan et al., 2025).

Among the detectors reviewed here, RevPRAG reports the highest true-positive rate (98%), above EmbedGuard's Tier-1 94.7% reference result. The numbers are not a head-to-head comparison because the systems use different evaluations. EmbedGuard's design addresses capabilities outside RevPRAG's scope, embedding provenance and cross-layer signal fusion, but the released package implements provenance only in software simulation and neither capability has an open head-to-head evaluation. A shared benchmark is required before ranking the systems directly.

Despite these advances, many evaluated defenses operate at one or two architectural levels. Recent concurrent work spans multiple stages, but shared-benchmark evidence for correlated four-layer decisions remains absent. Analysis of backdoor attacks on natural language generation provides insights into how adversaries embed backdoors at different abstraction levels—malicious training data provision, model parameter manipulation, and inference-time triggers (Sun et al., 2023). Studies demonstrate that data poisoning backdoors prove particularly challenging to detect as they exploit the model's learning process, typically assumed to be trustworthy.

2.4 Layered Defenses and Provenance Approaches

A wave of layered and provenance-oriented RAG defenses emerged in 2025-2026, and it is worth positioning EmbedGuard precisely against them rather than claiming the layered idea in isolation. Saleem et al. (2026) propose a three-layer framework against prompt injection in RAG chatbots — input anomaly classification, provenance-based context assembly, and output drift auditing — reporting a 27.3 percentage point improvement over the best single layer; the framework targets prompt injection specifically and evaluates layer complementarity rather than fusing signals into a joint decision. Ramakrishnan and Balaji (2025) combine embedding-based anomaly detection, prompt guardrails, and multi-stage response verification for agent security. Pallerla et al. (2026) orchestrate multiple defenses across attack vectors through a sentinel-strategist architecture, selecting defenses per detected vector rather than correlating their signals. In each case, layers act as sequential or orchestrated filters: a query passes or fails each stage independently. EmbedGuard's correlation engine instead combines per-layer scores into one decision (Section 3.6). With the package defaults, three layers each reporting a confidence-weighted score of 0.6 produce a fused score of 0.75 and a FLAG in gated mode, while two such layers produce 0.675 and do not. Whether this raises detection against attacks tuned below each layer's threshold has not been evaluated.

On the provenance axis, Wanger (2026) proposes software-signature provenance for vector stores, pinning each embedding to its source content and producing model via Ed25519 signatures. This validates embedding provenance as a defense primitive but leaves the signing keys and the embedding model itself inside the software trust boundary: an attacker who compromises the embedding host can sign poisoned vectors. Chrapek et al. (2024) demonstrate that full LLM inference pipelines can run inside Intel SGX/TDX enclaves with under 10% overhead, establishing the feasibility of TEE-hosted inference but without an embedding-specific attestation protocol. EmbedGuard's target design combines the two: embeddings would be generated inside a TEE whose attestation report binds the model hash and platform measurements, moving the signing root from software-held keys to hardware (Section 3.3). The released package implements only an HMAC simulation of this protocol. In either form, provenance detects vectors injected or modified outside the attested ingestion path; it does not detect poisoned documents that enter through that path (Section 4.3).

Post-PoisonedRAG detection research has also expanded at individual layers: gradient-masking detection (GMTP; Kim et al., 2025b), poisoning traceback for forensic attribution (RAGForensics; Zhang et al., 2025), LLM-activation-based detection (RevPRAG; Tan et al., 2025), adversarial-hubness detection in embedding space, and token-influence attribution. These methods strengthen individual layers and are complementary to EmbedGuard's architecture: any of them can serve as a drop-in signal source for the correlation engine.

2.5 Adaptive Adversaries and Threshold Probing

Query-efficient adversarial testing frameworks demonstrate how sophisticated adversaries optimize attacks against deployed defenses using Bayesian optimization methods, efficiently exploring attack spaces with low query budgets even against black-box defenses without internal knowledge (Lee et al., 2023). Adaptive attackers employ iterative processes that learn to optimize attacks through feedback from detection failures. Statistical threshold defenses prove particularly vulnerable as adversaries sample around threshold boundaries and design attacks exploiting these limits. Graph-theoretic approaches such as ReliabilityRAG provide provable guarantees under bounded-corruption assumptions. They do not address vectors written outside the ingestion path, which is the case hardware-backed attestation targets, and attestation in turn does not address poisoned text that enters through ingestion.

Table 2: Single-Layer Defense Limitations

Defense Mechanism	Primary Detection Target	Vulnerability to Adaptation	Evasion Strategy
Perplexity-Based Filtering	Linguistic anomalies in document content	High vulnerability to fluent text generation	Linguistically coherent malicious documents
Clustering-Based Outlier Detection	Spatial positioning in embedding space	Moderate vulnerability to centrality optimization	Embedding space centrality maintenance
Activation-Based Analysis	Model behavior during inference	Moderate vulnerability to normal pattern mimicry	Contextually appropriate activation patterns
Statistical Threshold Monitoring	Anomalous similarity distributions	High vulnerability to threshold probing	Systematic boundary identification

2.6 Geometric Properties Enabling Attacks

The mechanics of embedding-space attacks explain why conventional anomaly detection approaches prove insufficient for securing RAG systems. In high-dimensional embedding spaces, the curse of dimensionality creates regions unlikely to contain legitimate documents, providing exploitable opportunities for attackers. Adversaries can position documents near specific query vectors to secure preferential retrieval; whether such documents also evade distance-based outlier detection depends on the attack and on the reference distribution.

In high dimensions the contrast between nearest and farthest neighbor distances shrinks for many data distributions, which can weaken per-point distance tests, and an attacker who concentrates adversarial signal in query-relevant directions may exploit this. The effect is attack-specific: for PoisonedRAG passages, the distance between the retrieved-set mean and a clean baseline separated attack from clean retrievals (Section 4.3).

3. Materials and Methods

3.1 Architectural Overview

EmbedGuard implements a reference framework for reasoning about security signals across four RAG layers. The released code executes the enabled checks sequentially and contributes their signals to a deterministic correlation engine; the open benchmark measures only the pattern-based prompt path and does not establish full-pipeline production latency or throughput.




Figure 1: EmbedGuard cross-layer detection, traced through an illustrative attack. An attacker writes a fluent poisoned vector directly into the vector store, bypassing the attested ingestion path (red dashed path); a benign user query (blue path) then retrieves it. The prompt layer sees a clean query ($s_1 = 0.1$), the provenance layer finds no valid certificate for the out-of-band vector ($s_2 = 1.0$), retrieval analysis detects a shifted similarity distribution ($s_3 = 0.8$), and the output proxy observes answer instability under perturbation ($s_4 = 0.7$). With unit confidence for the illustration, the current package computes weighted consensus $1.325/1.80 = 0.736$, applies a strongest-signal floor of 1.000, adds the 0.150 three-layer correlation boost, and clips the final score to 1.000; active mode therefore returns BLOCK at the 0.85 threshold. The values are illustrative signals executed through the current fusion code, not measured layer outputs. Had the same document been submitted through the approved ingestion pipeline, it would carry a valid certificate and s_2 would be 0; detection would then depend on the retrieval, prompt, and output signals (Section 4.3).

Current implementation and evidence status:

Component	v1.2 implementation status	Open evidence status
Prompt detector	Implemented; 83-pattern path; neural mode is disabled by default and its fine-tuned checkpoint is not published	Exercised by the 135-sample Tier-2 regression benchmark
Embedding provenance	HMAC software simulator; target SEV-SNP protocol is design-only	Tamper-path unit tests; no hardware attestation evidence
Retrieval analyzer	Experimental PCA, Mahalanobis-distance, and temporal-rank prototype	Tier-2b PoisonedRAG evaluation (Section 4.3): package defaults detect 0/50; distance-only weights with a clean-calibrated threshold detect most attacks
Output verifier	Synthetic-output proxy; optional same-generator callback path	No dedicated open evaluation or deployed-LLM benchmark
Correlation engine	Implemented deterministic consensus/floor/boost algorithm	Unit-tested; no open four-layer ablation

Table 3: EmbedGuard Security Requirements

Requirement	Description	Verification Method
SR-1: Integrity	Detect poisoned or out-of-band vectors before retrieval influences generation	Provenance check for out-of-band vectors; retrieval-distribution score with a threshold calibrated on clean traffic (Section 4.3); fused four-layer decision not yet evaluated
SR-2: Provenance	Bind embeddings to an approved model identifier and a specific document hash	Attestation-certificate validation; software-simulated in this release
SR-3: Availability	Maintain system responsiveness under active attack conditions	Latency monitoring with <100ms target
SR-4: Auditability	Expose layer details for caller-owned incident records	Returned decision/layer details; log retention is caller-owned

3.2 Layer 1: Prompt Injection Detection

The released prompt layer performs lexical pattern analysis to identify selected injection attempts and jailbreak strings before input enters the retrieval pipeline. Recent research on universal adversarial attacks demonstrates systematic vulnerabilities in language model input processing, enabling adversaries to use specially crafted prompt suffixes to elicit malicious model outputs (Zou et al., 2023; Carlini et al., 2023). The implementation does not infer user intent or provide a trained semantic classifier by default.

The prompt analyzer employs a pattern-based classifier using 83 detection patterns covering diverse attack categories. It matches both the original query and a normalized representation that strips zero-width characters, applies Unicode NFKC normalization, and removes whitespace. Matched patterns contribute to a cumulative score using $\text{score} = \min(0.75 + (\text{num_matches} \times 0.05), 1.0)$, with a 0.70 decision threshold. Detection targets include direct instruction injection, jailbreak attempts, instruction smuggling, context manipulation, prompt leaking, role manipulation, indirect injection, encoding obfuscation (Unicode, Base64), delimiter confusion, XML/Markdown injection, hypothetical/fictional framing, translation attacks, authority claims, emotional manipulation, RAG-specific attacks, composite multi-vector attacks, and subtle manipulation. The current open benchmark observes 30/30 attacks detected and 0/105 benign queries flagged at 0.085ms aggregate mean latency on the recorded commodity-hardware run.

Detection signals from the prompt layer receive an intermediate default weight ($\text{beta}_1 = 0.35$) in the correlation engine because lexical patterns can match legitimate unusual queries. While prompt-layer detection prevents adversaries from using crafted queries to surface poisoned content, it provides insufficient protection against embedding-space poisoning, where legitimate queries unknowingly trigger the retrieval of malicious documents.

Table 4: Attack Pattern Taxonomy (83 patterns)

Category	Example signatures
Direct injection	ignore ... previous instructions, disregard ... prior
Role manipulation	you are now ..., pretend to be, act as
System extraction	show your prompt, reveal instructions
Delimiter attacks	[INST], < im_start >, ### System
Encoding bypass	base64, rot13, hexadecimal payloads
Jailbreak keywords	DAN, developer mode, bypass safety
Unicode and obfuscation	zero-width characters, repeated or split payloads
Context manipulation	priority claims, overrides, new directives
Framing attacks	hypothetical, fictional, translation, role-play
Social engineering	authority, emotional, and urgency claims
RAG-specific injection	retrieval-time instructions and document-context markers

Category	Example signatures
Composite patterns	multi-vector combinations

The full pattern taxonomy and regex definitions are available in `embedguard/prompt_detector/__init__.py` (INJECTION_PATTERNS).

3.3 Layer 2: Cryptographic Embedding Attestation

EmbedGuard's architectural proposal uses hardware-rooted attestation to bind an embedding to its source document, model, output vector, time, and execution platform. Trusted Execution Environments can isolate computation from privileged system software and produce platform evidence (AMD, 2020; Misono et al., 2024). The following is the target protocol described by the IJCESN version of record; it is not implemented by the repository's software simulator.

Target TEE-Based Embedding Generation Protocol:

In a hardware deployment, legitimate embeddings would be generated inside a measured SEV-SNP confidential VM:

- Confidential-Guest Initialization:** Embedding model (all-mpnet-base-v2, 768 dimensions) and source documents are loaded into encrypted guest memory protected from the hypervisor.
- Isolated Computation:** Vector generation executes in a hardware-isolated context. SEV-SNP supplies a launch measurement for the guest image; the protocol separately binds the approved model hash because dynamically loaded model weights are not automatically covered by that launch measurement.
- Report Binding:** The verifier issues a fresh challenge nonce N and records it in a replay cache. The confidential guest computes $E = f_Model(D)$ and a binding value $B = H(H(D) || H(Model) || H(E) || T || N)$, places B in the SEV-SNP report-data field, and requests an attestation report. The AMD Secure Processor signs a report that includes the report-data binding plus platform evidence such as launch measurement, policy, and TCB version. The bound metadata contains:
 - Input document hash: $H(D)$
 - Embedding model hash: $H(Model)$
 - Output vector hash: $H(E)$
 - Timestamp: T
 - Verifier challenge nonce: N
 - AMD SEV-SNP report fields such as launch measurement, policy, and TCB version


The platform attestation key is not an application signing key inside the guest; downstream verification relies on AMD's VCEK endorsement chain and the report-data binding.

Verification at Retrieval Time:

The target retrieval verifier would validate:

- Endorsement and Signature:** Validate the AMD ARK/ASK/VCEK chain, then verify the attestation-report signature with the validated VCEK
- Binding:** Recompute B from the document, approved model, embedding, timestamp, and challenge nonce and compare it with report data
- Freshness Policy:** Require an unused verifier-issued nonce, atomically mark it consumed, and reject reports outside the deployment's maximum-age window
- Platform Policy:** Check launch measurement, guest policy, and reported TCB against an allowlist

Target AMD SEV-SNP protocol — design only; not implemented in the released package



Released code uses software HMAC binding only; it obtains no SNP report or AMD endorsement chain.

Figure 2: Target TEE-based embedding attestation protocol. Documents are hashed and loaded into a measured SEV-SNP confidential VM, which generates embeddings and requests attestation evidence; retrieval verifies that evidence before accepting vectors. The released repository provides a software HMAC simulator of the binding and verification flow, not an AMD SEV-SNP attestation implementation.

Released Implementation Boundary:

EmbeddingAttestationLayer uses HMAC-SHA256 with a caller-supplied key or an ephemeral process key, a constant simulated model hash, document and embedding hashes, an ISO timestamp, and a validity period. Setting `tee_available=True` only adds clearly labeled simulated platform metadata; the code neither obtains an AMD SEV-SNP attestation report nor validates an endorsement chain. The simulator is useful for testing tamper detection and certificate plumbing, but it does not establish a hardware root of trust.

Security Properties:

Within the simulator, missing certificates are marked unverified and document-hash, embedding-hash, model-hash, expiry, or HMAC failures invalidate a certificate. A real TEE deployment could move the trust boundary to hardware, but the open package does not demonstrate that property. The default correlation weight remains $\beta_2 = 0.75$ because provenance failure is treated as the strongest layer signal. The layer's score is the fraction of retrieved documents without a valid certificate, so one uncertified vector among five retrieved documents yields a score of 0.2; per-document handling of certificate failures is left to the caller.

Performance Characteristics:

The IJCESEN version of record reports 1.8ms signature generation, 0.3ms validation per retrieved document, 2.1ms batch validation for 10 documents, and 12.8ms ingestion overhead. Those Tier-1 measurements are retained as publication history; the current repository does not include the hardware environment or a script that reproduces them.

3.4 Layer 3: Retrieval Distributional Analysis

The released retrieval analyzer combines three experimental signals. The Tier-2 prompt benchmark does not exercise this layer; the Tier-2b corpus-poisoning benchmark (Section 4.3) does. No threshold shipped with the package is calibrated for any corpus.

Incremental Principal Component Analysis:

IncrementalPCA retains a bounded window of embedding rows and recomputes a simplified SVD representation when its warm-up batch is reached and on scheduled updates. Until warm-up, the analyzer returns a zero PCA anomaly score with low confidence. After warm-up, reconstruction error is normalized to a [0,1] score. The default configuration retains 50 components, a 100-row fitting window, and schedules updates every 1,000 queries; the implementation is a research prototype rather than a production incremental-SVD algorithm.

Mathematical formulation:

- Given centered embedding matrix X in $\mathbb{R}^{(n \times d)}$ for n documents
- Let V_k contain the top- k right singular vectors of the fitting window
- Reconstruct: $X_{\text{hat}} = (X V_k^T) V_k$
- Anomaly score: $\text{mean}_i \|X_i - X_{\text{hat}_i}\|_2$, normalized by the implementation scale

Distribution-Distance Monitoring:

Despite the legacy internal name `_kl_divergence_score`, the released code computes a regularized Mahalanobis distance between the current mean embedding and an evolving baseline. The baseline is initialized from the first retrieved batch and, unless frozen, updated by an exponential moving average ($\alpha = 0.01$) after every scored query, including anomalous ones; the distance is scaled by $3/\sqrt{d}$ before clipping to [0,1]. It uses a diagonal approximation until the buffered rows (the last 100 retrieved batches) number at least $d + 10$, then applies a regularized pseudo-inverse of the full covariance. With top-5 retrieval and 768-dimensional embeddings the buffer holds at most 500 rows, so in that setting, including the evaluation of Section 4.3, only the diagonal path runs: a per-dimension scaled distance between the retrieved-set mean and the baseline mean, with the scales taken from a single early warm-up batch. The repository does not reproduce the previously stated $\tau = 0.15$, 89.1% detection, or 4.3% false-positive calibration claims, so they are not presented as current results.

Temporal Rank Correlation:

After at least 10 historical observations, the analyzer compares the current retrieval's document ranking with the most recent earlier query whose retrieved document-ID set shares at least two documents. SciPy's `spearmanr` computes Pearson correlation over the tie-aware average ranks of those shared documents only:

$\rho = \text{corr}(\text{rank_avg}(\text{pos}_t(D)), \text{rank_avg}(\text{pos}_j(D)))$, D = shared document IDs

If correlation is undefined for constant inputs, the implementation maps ρ to 0. The rank anomaly is $\max(0, (1 - \rho) / 2)$; when ρ is below the configured 0.30 minimum, it adds 0.30 and caps the result at 1.0. When no earlier query is comparable, the component abstains with confidence 0 and is excluded from fusion.

Manuscript v3.1 and package versions up to 1.2.0 instead compared the current query's retrieval-score vector with the preceding query's vector, right-padded with zeros. Both vectors are sorted in descending order, so their ranks always agree, ρ is 1, and the component returned 0 for every query while carrying weight 0.2 in the fusion. Version 1.3.0

replaces that comparison with the shared-document ranking above.

Component fusion and calibration. The analyzer weights each component by its configured weight (default 0.5, 0.3 and 0.2 for PCA, distribution distance and rank) multiplied by the component's confidence, excludes components with zero weight or zero confidence, and renormalises. The weights are configurable. `calibrate_threshold` returns the (1 - target FPR) quantile of fused scores on held-out clean traffic, and `freeze_baseline` stops scored requests from updating the PCA model, covariance buffer and baseline mean. Without freezing, a sustained run of shifted retrievals enters the covariance estimate it is scored against and the distance collapses; the unit test `test_adaptive_baseline_is_desensitised_by_a_sustained_shift` reproduces this. It does not condition on user-session semantics. Retrieval receives $\beta_3 = 0.50$ in cross-layer fusion. Its behaviour against published corpus-poisoning texts is measured in Section 4.3.

3.5 Layer 4: Output Consistency Verification

The output layer is a development proxy for perturbation-based stability testing, motivated by work showing that backdoored models can be sensitive to input perturbations (Sun et al., 2023). It is not an end-to-end LLM evaluation in the released repository.

Retrieval Perturbation Strategy:

The implementation generates $K=5$ document-set perturbations by cycling through document removal, document reordering, sentence/chunk removal, and small text-noise insertion.

Stability Measurement:

By default, the implementation constructs deterministic synthetic text from document content for both the baseline and perturbed sets. For real-output testing, the caller must supply an `output_generator` callback that reruns the same generator on the original and every perturbed document set; an optional `generated_output` may serve as the baseline only when that callback is present. A real output supplied without generator access is marked `not_evaluated`, never compared against synthetic proxies. It then measures output stability:

$$\text{Stability} = (1/K) * \sum_{i=1}^K \text{sim}(\text{output_original}, \text{output}_i)$$

When the optional semantic model is explicitly enabled, `sim()` is cosine similarity in embedding space; the default path uses deterministic Jaccard word overlap and performs no model download. The default stability threshold is 0.65 and the returned anomaly score is one minus mean similarity. The open benchmark does not validate benign/attack separation, trigger frequency, or latency for this layer. Its lower fusion weight, $\beta_4 = 0.20$, reflects the risk that ordinary output variation can look unstable.

3.6 Threat Correlation Engine

The released correlation engine is deterministic; it does not contain a learned classifier or Bayesian posterior. For layer score s_i , confidence c_i , and default weight β_i , it computes a confidence-weighted consensus and a fail-closed strongest-signal floor:

$$\text{consensus} = (\sum_i \beta_i c_i s_i) / (\sum_i \beta_i c_i)$$

$$\text{base} = \max(\text{consensus}, \max_i(c_i s_i))$$

It then counts confidence-bearing elevated layers with $c_i s_i > 0.5$. For k elevated layers, the additive correlation factor is zero when $k \leq 1$ and $0.075(k-1)$ otherwise, using the package default correlation boost of 0.15. Zero-confidence scores cannot contribute to this boost. The final score is $\min(1, \text{base} + \text{correlation factor})$. Gated mode flags at a score of at least 0.70; active mode blocks at a score of at least 0.85 and flags scores from 0.70 to below 0.85.

Table 5: Default Layer Weights in the Released Correlation Engine

Layer	Weight (beta)	Rationale
Prompt	0.35	Pattern and optional neural detection can produce false alarms
Embedding provenance (HMAC simulator; TEE target)	0.75	Provenance failure is the strongest trust signal
Retrieval	0.50	Distributional analysis is informative but statistical
Output	0.20	Legitimate generation variance can cause instability

These defaults are implementation parameters, not a newly reproduced calibration result. The repository does not release the claimed 5,000-sample calibration set.

Cross-Layer Pattern Recognition:

The engine adds a fixed increment when two or more layers report confidence-weighted scores above 0.5. The illustrative case in Figure 1 is:

- a fluent document that does not trigger the prompt layer;
- a vector written outside the attested ingestion path, so it carries no valid certificate;
- a shift in the retrieved-embedding distribution;
- instability of the output proxy under perturbation.

For poisoned text admitted through the normal ingestion pipeline the second signal is absent (Section 4.3). Whether sub-threshold signals on several layers combine to raise detection has not been measured.

For the illustrative unit-confidence signals in Figure 1, the raw weighted sum is 1.325 and the weight total is 1.80, so consensus is 0.736. The strongest-signal floor is 1.000, three layers are elevated, and the 0.150 correlation factor is added before clipping; the package therefore returns a final score of 1.000 and BLOCK in active mode. This example demonstrates the code path rather than an empirical calibration result.

3.7 Operational Modes

EmbedGuard supports three deployment modes accommodating diverse operational requirements:

Passive Mode: The correlation engine always returns LOG, irrespective of score. This supports observation without automated intervention; downstream storage, retention, and incident-correlation behavior are deployment responsibilities rather than benchmarked package features.

Gated Mode: The engine returns FLAG at scores of at least 0.70 and ALLOW below 0.70. The package returns the decision and layer details; pausing a request, presenting a review interface, and measuring analyst review time require an external integration and are not evaluated here.

Active Mode: The engine returns BLOCK at scores of at least 0.85, FLAG from 0.70 to below 0.85, and ALLOW below 0.70. The package does not itself generate a fallback response; the calling application owns that policy and must calibrate thresholds against its own traffic before enabling automatic blocking.

4. Results and Discussion

We report results at two evaluation tiers with different reproducibility properties, and we are explicit about which claims each tier supports.

Tier 1 — Archived version-of-record claims. The IJCESEN article reports a production-scale four-layer evaluation on 500,000 embeddings and 47,000 queries with AMD SEV-SNP hardware. The repository does not contain the raw predictions, production corpus, attack generator, baseline configurations, or attestation logs required to audit that experiment. Tier-1 values below are preserved to document the immutable published record (DOI 10.22399/ijcesen.4869); they are not independently reproducible from this artifact and do not validate the repaired post-publication implementation.

Tier 2 — Open regression benchmark (repeatable on commodity hardware). The EmbedGuard open benchmark contains 135 locally curated samples across three benign benchmark-style files and a 25-category injection file, exercising the released prompt-layer detector in pattern-only mode without specialized hardware. This tier is repeatable from the released repository with a single command and supports only the observed regression-set prompt-detection, specificity, and latency claims. It does not evidence generalization, cross-layer benefit, TEE behavior, or comparative Tier-1 claims.

4.1 Archived Tier-1 Version-of-Record Results (Not Reproduced Here)

Published protocol description. The version of record states that four attack families were used (optimization-based, transferability-based, semantic manipulation, and adaptive attacks), but its results table reports a fifth row, Coordinated Multi-Layer, whose construction it does not describe. Because the corresponding generator, configurations, and per-example outputs are unavailable here, the repository cannot verify the stated learning rate, iteration count, transfer rate, adaptive optimization procedure, or the composition of the fifth row.

Table 6a: Archived Tier-1 Detection Performance Reported in the Version of Record

Attack Type	Detection Rate	False Positive Rate	Mean Latency
Optimization-Based	94.7%	3.2%	47ms
Transferability-Based	91.4%	4.1%	51ms
Semantic Manipulation	88.9%	3.8%	49ms
Adaptive Attacks	89.3%	5.2%	53ms
Coordinated Multi-Layer	96.2%	2.9%	58ms

The version of record reports its highest detection for coordinated multi-layer attacks (96.2%) and 89.3% for adaptive attacks. Without raw predictions and attack-generation artifacts, this repository cannot verify the proposed explanation that correlation caused the difference.

Table 6b: Archived Tier-1 Comparative Values Reported in the Version of Record

Defense System	Baseline Detection	Adaptive Detection	FPR	Latency
EmbedGuard	94.7%	89.3%	3.2%	51ms
RAGuard	87.2%	61.4%	4.8%	38ms
RobustRAG	82.9%	58.7%	6.1%	42ms
TrustRAG	79.3%	54.2%	5.3%	35ms

As published, the EmbedGuard row combines the Optimization-Based detection rate and false-positive rate from Table 6a (94.7%, 3.2%) with the latency averaged over all attack rows (51 ms). The per-category sample counts in the version of record (12,500, 9,800, 11,200, 8,300 and 5,200) sum to its entire 47,000-query workload, and it describes no separate benign query set, so the denominator of every Tier-1 false-positive rate is unstated. Weighted by those counts, detection is 91.8% and the false-positive rate is 3.85%.

The published table reports a 27.9–35.1 percentage-point adaptive-detection difference and an 18.4 percentage-point full-system versus embedding-only ablation difference. Those arithmetic differences are retained as publication history, not as independently auditable comparative or causal evidence in this repository.

4.2 Tier 2: Open Benchmark Evaluation

4.2.1 Experimental Setup

Infrastructure Configuration:

- Recorded host: macOS arm64 (Apple silicon), Python 3.14.3, embedguard 1.2.0, source commit 2bc75a1; the package supports Python 3.10+
- Detector: production PromptInjectionDetector in pattern-only mode, 83 patterns, threshold 0.70
- Recorded run: 2026-07-10; committed JSON and Markdown report under results/

Benchmark Datasets:

The repository contains locally curated benchmark-style benign queries labeled after three public NLP datasets plus a curated injection regression set. It does not include upstream example identifiers, extraction manifests, or checksums that would establish these rows as exact subsets of those datasets:

Dataset	Samples	Type	Source
Natural Questions-style	50	Curated benign queries	Local file labeled NQ
HotpotQA-style	25	Curated benign queries	Local file labeled HotpotQA
MS-MARCO-style	25	Curated benign queries	Local file labeled MS-MARCO
Injection regression set	35	30 attacks + 5 benign	Local 25-category set
Total	135	—	—

Attack Dataset Composition:

The injection attack dataset spans 25 distinct attack categories derived from recent security literature (Zou et al., 2025; Liu et al., 2023; Carlini et al., 2023):

Attack Category	Samples	Description
Direct Instruction	2	Explicit command injection
Jailbreak Attempt	1	Safety bypass attempts
Instruction Smuggling	1	Hidden instruction embedding
Context Manipulation	1	Context window exploitation
Prompt Leaking	1	System prompt extraction
Role Manipulation	1	Role/persona hijacking
Indirect Injection	1	Third-party content injection
Unicode Obfuscation	1	Unicode character exploitation
Base64 Encoding	1	Encoded payload delivery
Delimiter Confusion	1	Structural boundary attacks

Attack Category	Samples	Description
XML Injection	1	XML/markup exploitation
Markdown Injection	1	Markdown rendering attacks
Developer Mode	1	Privileged mode claims
Hypothetical Framing	1	Hypothetical scenario abuse
Fictional Framing	1	Fiction context exploitation
Translation Attack	1	Multi-language evasion
Repetition Attack	1	Repetition-based injection
Authority Claim	1	False authority assertion
Emotional Manipulation	1	Emotional appeal exploitation
RAG-Specific	3	Retrieval-targeted attacks
Composite Attack	2	Multi-vector combinations
Subtle Manipulation	2	Low-signal attacks
Multi-Turn Setup	1	Conversation state exploitation
Payload Splitting	1	Fragmented payload delivery
Virtualization	1	Virtual environment claims
Total Attacks	30	—
Benign Control	5	Legitimate queries

Reproducibility: All Tier-2 datasets, detector code, and evaluation scripts are released at <https://github.com/neerazz/embedguard>. Classification counts are deterministic for the committed inputs; wall-clock latency varies by hardware and system load.

4.2.2 Detection Performance Results

Table 7: EmbedGuard Aggregate Detection Performance

Metric	Value	Description
Total Samples	135	Across all benchmarks
Attack Samples	30	Malicious queries
Benign Samples	105	Legitimate queries (5 control + 100 benchmark)
True Positives	30	Attacks correctly detected
False Positives	0	Benign queries incorrectly flagged
True Negatives	105	Benign queries correctly passed
False Negatives	0	Attacks missed
Detection Rate	100%	30/30 attacks detected
False Positive Rate	0%	0/105 benign flagged
Accuracy	100%	(30+105)/135
Precision	100%	30/(30+0)
Recall	100%	30/(30+0)
F1 Score	1.00	Harmonic mean

Table 8: Detection Performance by Attack Category

Attack Category	Samples	Detected	Detection Rate	Mean Score
Direct Instruction	2	2	100%	0.95
Jailbreak Attempt	1	1	100%	1.00
Instruction Smuggling	1	1	100%	0.80
Context Manipulation	1	1	100%	0.85
Prompt Leaking	1	1	100%	0.85
Role Manipulation	1	1	100%	0.85
Indirect Injection	1	1	100%	0.85
Unicode Obfuscation	1	1	100%	0.85
Base64 Encoding	1	1	100%	0.85
Delimiter Confusion	1	1	100%	0.80
XML Injection	1	1	100%	0.90
Markdown Injection	1	1	100%	0.80
Developer Mode	1	1	100%	0.90
Hypothetical Framing	1	1	100%	0.90
Fictional Framing	1	1	100%	0.80
Translation Attack	1	1	100%	0.90
Repetition Attack	1	1	100%	0.85
Authority Claim	1	1	100%	0.90
Emotional Manipulation	1	1	100%	0.90
RAG-Specific	3	3	100%	0.88
Composite Attack	2	2	100%	0.90
Subtle Manipulation	2	2	100%	0.93
Multi-Turn Setup	1	1	100%	0.90
Payload Splitting	1	1	100%	0.85
Virtualization	1	1	100%	0.95

On these fixed samples, the 83-pattern detector identifies all attacks and flags none of the benign queries. This is an observed benchmark result, not a guarantee for unseen attacks or production traffic. Attack-category sample sizes range from one to three, so the per-category percentages are descriptive rather than stable population estimates.

Table 9: Benign Query Performance by Benchmark Dataset

Dataset	Samples	False Positives	True Negatives	Specificity	Mean Latency
Natural Questions-style	50	0	50	100%	0.042ms
HotpotQA-style	25	0	25	100%	0.099ms
MS-MARCO-style	25	0	25	100%	0.038ms
Control (in injection set)	5	0	5	100%	—
Total Benign	105	0	105	100%	—

The detector flags none of the 105 benign strings in these three fixed local files. Because the samples lack upstream provenance and were not selected as hard negatives, this result is a regression observation rather than evidence of robust specificity across legitimate-query distributions.

4.2.3 Uncertainty Analysis

We report two-sided 95% Wilson score intervals over the observed confusion-matrix counts. We do not apply signed-rank, McNemar, or effect-size tests because this benchmark has no paired comparator and the released result summary does not contain an independent distribution of per-sample model scores suitable for those analyses.

Quantity	Observed count	Point estimate	95% Wilson interval
Detection rate	30/30	100%	88.6%-100%
Specificity	105/105	100%	96.5%-100%

The intervals are generated by `scripts/statistical_tests.py` from the committed benchmark JSON. The 88.6% lower bound for detection makes the central limitation visible: 30 successful attack detections do not establish universal coverage, even though the point estimate on this fixed benchmark is 100%.


4.2.4 Latency Analysis

Table 10: Per-Dataset Latency Statistics

Dataset	Mean (ms)	Median (ms)	P95 (ms)	P99 (ms)	Min (ms)	Max (ms)
Injection set (30 attacks + 5 benign)	0.169	0.179	0.249	0.273	0.045	0.275
Natural Questions-style	0.042	0.038	0.072	0.087	0.017	0.095
HotpotQA-style	0.099	0.092	0.181	0.187	0.055	0.187
MS-MARCO-style	0.038	0.033	0.073	0.083	0.016	0.085
Aggregate	0.085	0.055	0.219	0.259	0.016	0.275

Recorded on 2026-07-10 using the released pattern-only detector (83 patterns, threshold 0.70, N=135). Timing is machine- and load-dependent; classification counts are the reproducible result.

Latency is sub-millisecond on this recorded run, with an aggregate mean of 0.085ms and P99 of 0.259ms. The benchmark is single-threaded and does not measure service-level throughput, concurrency, tail behavior under load, or end-to-end RAG latency.



Source: results/benchmark_results_20260710_025640.json · host-dependent timing

Figure 3: Single-run, host-dependent prompt-detector latency for the committed 135-sample Tier-2 run. Circles are means, diamonds are p95, triangles are p99, and gray whiskers span minimum to maximum. Dataset names denote locally curated benchmark-style inputs, not verified upstream subsets.

4.2.5 Evaluation Boundary

The open prompt benchmark does not measure concurrent throughput, full-pipeline RAG latency, semantic or model-based prompt detection, adaptive attacks, or performance on JailbreakBench (Chao et al., 2024). Corpus poisoning is evaluated separately, for the provenance and retrieval layers and one published attack, in Section 4.3.

4.3 Tier 2b: Open Corpus-Poisoning Evaluation (PoisonedRAG)

The Tier-2 benchmark exercises the prompt layer, while the attack this paper is concerned with arrives through the corpus: a benign query retrieves passages that an attacker planted. Tier 2b tests that case with published attack texts and a public corpus. It evaluates the provenance and retrieval layers. In the end-to-end runs only the retrieval layer is enabled: the queries are benign, so the prompt layer has nothing to detect; there is no generator for the output proxy; and provenance is scored separately.

4.3.1 Protocol

Attack and corpus. The attack texts are the black-box adversarial passages released with PoisonedRAG for Natural Questions (Zou et al., 2025; Geng et al., 2024): five LLM-generated passages for each of 100 target questions, each written to make an LLM give an attacker-chosen answer. The *black-box* form prepends the target question to each passage, as in the released attack code. The *stealth* form uses the passage alone, which removes the verbatim copy of the question that makes black-box poison lexically close to its target. The clean corpus is BEIR Natural Questions (Kwiatkowski et al., 2019; Thakur et al., 2021). Retrieval uses sentence-transformers/all-mpnet-base-v2 (Reimers

and Gurevych, 2019) with normalised cosine similarity and top-5 results over a pool of 5,658 clean passages: the relevance-judged passages of every evaluated query plus 5,000 passages drawn at random from the 2,681,468-passage corpus. One, three or five poison passages per target question are added to the pool. All four input files are fetched from pinned upstream revisions and checked against SHA-256 digests by `scripts/fetch_tier2b_data.py`.

Splits and calibration. A seeded split (seed 12) divides the 100 target questions into 50 for calibration and 50 for test, and draws 450 non-target NQ queries into warm-up (150), calibration (100) and test (200) sets. The detector sees only clean traffic before testing: `EmbedGuard.analyze()` is run over the 150 warm-up queries, then `EmbedGuard.calibrate()` runs the 100 calibration-clean queries, sets the FLAG threshold to the clean-score quantile for the target false-positive rate, and freezes the retrieval baseline. No attack example is used to fit the threshold or the baseline. The test stream then mixes the 50 test attacks with the 200 test-clean queries in random order. A query counts as detected when the decision is FLAG or BLOCK.

Retrieval weights. The end-to-end runs use `retrieval_component_weights = {pca: 0, kl: 1, rank: 0}`, that is, the distribution-distance component alone, which in this setting is the diagonal distance described in Section 3.4. A selection rule applied to the calibration split (keep components whose calibration-split AUROC is at least 0.6) selects the distance component alone in all six configurations; its calibration-split AUROC ranges from 0.74 (one poison passage per target) to 0.998 (five), while PCA ranges from 0.50 to 0.57 and rank from 0.42 to 0.46. This rule uses the labels of the 50 calibration attacks, and both the distance-only setting and the 0.6 cut-off were chosen after a pilot run on this dataset with the same split procedure. The test split was not used to choose them, but the choice is not independent of this benchmark. The per-component AUROCs reported here and in Section 4.3.3 come from a separate diagnostic analyzer that uses the package defaults and is frozen after the same 150-query warm-up; its rank component still reads query history and so depends on stream order.

Controls. Two controls separate a response to the poison from a response to the question. The first is the 200 test-clean queries. The second runs the same 50 test target questions against the clean pool, without poison; if these were flagged at the attack rate, the detector would be reacting to the questions rather than to the planted passages. Provenance is scored under two threat models: poison written into the store out of band, with no certificate, and poison submitted through the approved ingestion pipeline, which issues it a valid HMAC certificate.

4.3.2 Results

Table 11: Tier-2b test split, `EmbedGuard.analyze()` after `EmbedGuard.calibrate()`, 50 attacks and 200 clean queries per row

Poison form	Poison passages per target	Attack retrieval success	Detected at 5% target FPR	Detected at 1% target FPR	AUROC
Black-box	5	50/50	50/50 (92.9–100%)	50/50 (92.9–100%)	1.000
Black-box	3	50/50	43/50 (73.8–93.0%)	18/50 (24.1–49.9%)	0.985
Black-box	1	50/50	7/50 (7.0–26.2%)	0/50 (0.0–7.1%)	0.669
Stealth	5	50/50	48/50 (86.5–98.9%)	35/50 (56.2–80.9%)	0.996
Stealth	3	50/50	33/50 (52.2–77.6%)	11/50 (12.8–35.2%)	0.961
Stealth	1	50/50	7/50 (7.0–26.2%)	0/50 (0.0–7.1%)	0.658

Intervals are two-sided 95% Wilson intervals (Wilson, 1927). In every row the 200 test-clean queries were flagged 6 times at the 5% target (3.0%, 1.4–6.4%) and once at the 1% target (0.5%, 0.1–2.8%). The target questions against the clean pool were flagged 3/50 times at the 5% target and 0/50 at the 1% target, close to the clean-query rate and far below the attack rate at three and five passages, so the detections respond to the poison rather than to the questions. Attack retrieval succeeded in all 50 test cases in every configuration, with on average 1.1, 3.0, and 4.3 (stealth) to 4.8 (black-box) poison passages in the top five for one, three, and five injected passages.

Table 12: Other configurations on the same test split (black-box form, five passages per target)

Configuration	Attacks detected	Clean flagged
Retrieval analyzer, package-default weights (0.5/0.3/0.2) and 0.5 cut	0/50	0/200
Provenance, poison written out of band without a certificate	50/50	0/200
Provenance, poison ingested through the approved pipeline	0/50	0/200

4.3.3 Interpretation

Provenance does not address this attack. Attestation shows which model embedded which bytes; it does not show that the bytes are benign. Every PoisonedRAG passage submitted through the approved pipeline received a valid certificate and was accepted in all six configurations. Provenance caught every passage written into the store out of band, which is the attack it is designed for. The Tier-1 ablation (Section 4.1) attributes its largest single effect to the embedding TEE; the version of record does not state whether its poisoned documents entered through the attested ingestion path, and if they did, that attribution is inconsistent with this result.

The retrieval layer detects concentrated poisoning, not a single planted passage. The distance component responds to the mean embedding of the retrieved set moving away from the clean baseline. With three or five of the top five results drawn from poison, that shift is large enough to separate attack from clean retrievals (AUROC 0.96 to 1.00). With one planted passage per target, the mean moves little: AUROC falls to 0.66–0.67, the detector catches 7 of 50 at a 3.0% clean flag rate, and none at the 1% target. PoisonedRAG reports that attack success rises with the number of injected passages up to the retrieval depth (Zou et al., 2025, Section 5.3), so the regime in which this detector is blind is also the one in which the attack is weakest, although their Figure 4 shows that attack success with a single injected passage is not negligible. Retrieval of a planted passage is not the same as a successful attack, and this evaluation does not measure end-to-end attack success; a one-passages attack that does succeed would not be caught.

The package defaults do not detect the attack. With the default component weights and the uncalibrated 0.5 cut the retrieval layer detected 0/50. The PCA reconstruction-error component and the rank component carry no signal in the useful direction on this workload (test AUROC 0.46–0.53 and 0.35–0.49 for PCA and rank), yet they hold 70% of the default weight. Package version 1.3.0 therefore adds `retrieval_component_weights`, `EmbedGuard.calibrate()` and baseline freezing, and keeps the previous defaults unchanged for compatibility; the configuration evaluated here must be selected explicitly.

Scope. The pool is 0.2% of the NQ corpus, which makes the attacker's retrieval easier and the retrieved set more homogeneous than in full-corpus retrieval; the effect on detection in either direction is not measured here. The attacker is non-adaptive: an attacker who knows a distance detector is present could craft passages that stay close to the clean distribution, and no such attack is evaluated. The test split contains 50 attacks, so a rate of 50/50 is compatible with a true rate as low as 92.9%. The result covers one dataset, one embedding model and one attack family, and it neither reproduces nor supports the Tier-1 numbers in Section 4.1.

Reproduction. `python scripts/fetch_tier2b_data.py --data-dir data/tier2b` downloads and verifies the inputs (about 765 MB), and `scripts/run_tier2b.sh data/tier2b` regenerates all six configurations as `results/tier2b_poisonedrag_<form>_k<n>_20260928.{json,md}`, each recording input digests, library versions, the source commit and per-query rows. The committed files were generated from a clean checkout of code commit 80873d7, which contains every code change in this release; the only uncommitted files each run lists are the result files of the runs before it. One configuration takes about one minute on an Apple M-series GPU.

4.4 Architectural Analysis

Reconciling the two tiers:

The prompt classifier achieves 100% detection on the Tier 2 benchmark, while the full framework reports 89.3-96.2% on Tier 1. These are not contradictory; they measure different things. Tier 2 consists of prompt-borne attacks delivered at benchmark scale where the 83-pattern taxonomy covers every included sample. Tier 1's attack families (gradient-optimized document poisoning, cross-model transfer, adaptive evasion) largely bypass the prompt layer: the malicious content arrives through the corpus, not the query. A reader should treat Tier 2 as reproducible evidence for one prompt-layer implementation and Tier 1 as the published record of the broader architecture, with one caution: the Tier-1 ablation attributes its largest effect to attestation, but the version of record does not say whether its poisoned documents entered through the attested ingestion path. Poison that did would have received valid certificates (Section 4.3):

Layer	Detection Mechanism	Primary Target	Weight (β)
Prompt	Pattern matching (83 patterns)	Injection attempts, jailbreaks	0.35
Embedding	Target SEV-SNP attestation; repository HMAC simulator	Unauthenticated embeddings	0.75
Retrieval	PCA reconstruction error and Mahalanobis distance	Coordinated poisoning	0.50
Output	Perturbation stability testing	Hidden backdoors	0.20

The current Threat Correlation Engine computes a confidence-weighted consensus, takes the maximum of that consensus and the strongest confidence-adjusted layer signal, adds a bounded multi-layer correlation factor, and clips the result to 1.0. Gated mode flags at 0.70; active mode blocks at 0.85. Tier 2 supports only the prompt-layer observation. The archived Tier-1 results report broader-layer performance but cannot be reproduced from the open artifact.




Figure 4: Machine-readable transcription of the version-of-record Tier-1 ablation table (paper/data/tier1_ablation_vor.json). The 18.4-point difference compares the full system with the only single-layer configuration the version of record reports (embedding only), which it labels the best single layer; it is historical

comparative evidence, not causal isolation of the correlation mechanism. Tier-1 sample counts, confidence intervals, and raw observations are unavailable in the open archive.

4.4.1 What the Open Benchmarks Do Not Establish

The released Tier-2 data do not support a pattern-count ablation, a neural-versus-pattern comparison, or a cross-layer ablation because only the pattern-only prompt path is exercised. Earlier extended drafts contained estimated values for those comparisons; this revision removes them. The only cross-layer ablation retained is the Tier-1 table reported in the published version of record and plotted in Figure 4. Section 4.3 evaluates the provenance and retrieval layers separately on one corpus-poisoning workload; no open evaluation yet runs all four layers and the fused decision on the same workload, so gains from defense in depth remain unmeasured.

4.5 Limitations

This study has several limitations that warrant discussion:

1. **TEE and Source-Admission Assumptions:** The target attestation design assumes AMD SEV-SNP trust roots are not compromised and firmware is current. Recent vulnerabilities demonstrate ongoing security challenges: CVE-2024-56161 (CVSS 7.2) enables malicious microcode injection through improper signature verification, potentially allowing attackers with local administrator access to compromise confidential guests (AMD, 2025a). CVE-2024-21944 enables memory aliasing attacks that can undermine SEV-SNP integrity features (AMD, 2025b). These vulnerabilities require timely firmware patches. More fundamentally, provenance is not semantic trust: an authorized but malicious document can be faithfully embedded and attested. Source admission, authorization, and content validation remain separate controls. The released HMAC simulator provides no hardware guarantee.
2. **Benchmark Scale:** The open regression evaluation comprises 135 locally curated strings across one injection file and three benign style subsets. It measures fixed-file behavior, not generalization to the upstream Natural Questions, HotpotQA, or MS MARCO datasets. The released files and runner permit extension, but production traffic may contain distributions and attacks absent from this set.
3. **Perfect Detection Caveat:** The 100% detection rate on this benchmark reflects coverage of the 30 included attack strings. Novel attacks not covered by the 83 patterns may evade the prompt layer; the Wilson lower bound of 88.6% quantifies part of the sampling uncertainty but does not cover distribution shift.
4. **English-Language Focus:** Evaluation focused on English-language corpora. Multilingual RAG systems may exhibit different attack surfaces (e.g., Unicode exploitation across writing systems, translation-based evasion) not fully captured in our analysis.
5. **Adaptive Adversary Evolution:** Adversaries aware of EmbedGuard's pattern-based detection may develop novel evasion techniques exploiting pattern gaps. A complete implementation of the cross-layer architecture could provide redundancy, but the open benchmark does not show that the current provenance, retrieval, or output prototypes detect those evasions.
6. **Single-Vector Attack Concentration:** A complete cross-layer implementation could still miss attacks concentrated on one architectural layer. Correct provenance is not semantic trust: an authorized but malicious source could generate a correctly attested embedding while targeting retrieval or output behavior. Section 4.3 tests this case for the provenance layer: PoisonedRAG passages ingested through the approved pipeline received valid certificates and none was flagged (0/50).

7. **Attack Scope:** This work focuses on integrity attacks (content manipulation through injection and poisoning). Availability attacks such as jamming and denial-of-service are a complementary threat model outside this implementation and evaluation. Future work could extend EmbedGuard to detect availability violations through retrieval denial patterns.
8. **Concurrent Work:** Layered RAG defense and embedding provenance are active research areas. Earlier work combines stages of RAG defense (Xiang et al., 2024; Yao et al., 2025; Ramakrishnan and Balaji, 2025), applies cryptographic integrity to RAG content (Zhou et al., 2025b), and runs retrieval-augmented or foundation-model inference in confidential-computing environments (Chrapek et al., 2024; Addison et al., 2024); later frameworks (Section 2.4) pursue related designs. The combination proposed here, per-embedding attestation that binds document, model, and vector hashes in an SEV-SNP report together with fusion of prompt, retrieval, and output signals, was not found in the literature we searched. It remains a target design, and individual components have close relatives that practitioners may prefer for specific deployments — e.g., software-signature provenance where TEE hardware is unavailable. A head-to-head comparison against post-2025 detectors (GMTP, RevPRAG, hubness-based methods) on a shared benchmark is the highest-value next experiment.

4.6 Failure Mode Analysis

EmbedGuard's detection capabilities have theoretical and empirical limits that we document for transparency:

4.6.1 Pattern Evasion Techniques

Attackers can bypass pattern-based detection through several techniques:

Evasion Technique	Example	Detection Status	Mitigation
Whitespace injection	"I g n o r e"	Covered by current tests	Whitespace-stripped parallel match
Homoglyph substitution	Compatibility-equivalent Unicode forms	Partially covered	Unicode NFKC normalization
Synonym substitution	"Discard" for "Ignore"	⚠️ Vulnerable	Semantic embedding layer
Tokenization attacks	Unusual byte sequences	⚠️ Vulnerable	Byte-level pattern matching

Mitigation Implementation: The token normalization preprocessor (Section 3.2) removes whitespace, normalizes Unicode via NFKC, and strips zero-width characters before pattern matching. It improves coverage for the tested whitespace and invisible-character evasions, but NFKC does not map every cross-script homoglyph; synonym, transliteration, and byte-level evasions remain open risks.

4.6.2 TEE Compromise Scenarios

The attestation layer assumes uncompromised TEE firmware:

Vulnerability	CVE	CVSS	Attack Requirements	Impact on EmbedGuard
Microcode injection	CVE-2024-56161	7.2	Local admin access	Attestation bypass
Memory aliasing	CVE-2024-21944	6.5	VM guest privileges	Integrity violation

Threat Model Boundary: We explicitly exclude: (1) Physical attacks requiring hardware access; (2) insiders who control the approved ingestion pipeline or its source-admission policy, whose documents are attested by design (Section 4.5, item 1). Host administrators are in scope: they are the adversary hardware attestation targets, and CVE-2024-56161 shows that an unpatched platform lets such an adversary bypass it; (3) Nation-state adversaries capable of silicon-level compromise; (4) Supply chain attacks on TEE firmware distribution.

4.6.3 Distribution Shift Vulnerabilities

The 135-sample benchmark may not represent all production distributions:

Shift Type	Example	Detection Impact	Mitigation
Novel attack classes	Future jailbreak variants	Potential false negatives	Pattern updates; cross-layer redundancy (unevaluated)
Query length extremes	10K+ token queries	Untested	Length-aware thresholds
Domain terminology	Medical/legal jargon	Potential false positives	Domain-specific tuning
Multilingual content	Non-English attacks	Reduced coverage	Multilingual pattern extension

Recommended Deployment Practice: Organizations should fine-tune detection thresholds on domain-specific query samples before production deployment.

5. Applications and Societal Implications

RAG system integrity is critical for several application domains where safety or compliance with regulations is essential.

Healthcare: Clinical decision support systems often use RAG architectures to analyze medical literature, treatment protocols, clinical cases, and drug databases to produce evidence-based recommendations. Breaching the integrity of such systems has consequences for patient safety, regulatory compliance, and financial costs (IBM Security, 2024). A production implementation of the target protocol could prove which document and model produced an embedding; it could not prove that the medical source or generated recommendation is correct.

Financial Services: Financial institutions use RAG systems for trading, risk assessment, and regulatory compliance. Attacks on financial AI systems could create significant profit opportunities for adversaries or harm market integrity. EmbedGuard is a reference design for evaluating integrity controls in that setting, not a validated financial-services control.

Legal Research: Legal RAG systems retrieve case law, statutes, regulations, and legal commentary to inform legal analysis and brief generation. Compromised systems can introduce incorrect interpretation into client matters. Provenance can bind retrieved vectors to specified source bytes, but legal source authority and output correctness still require independent validation.

Equity Considerations: More than half of the breached organizations in IBM's 2024 study reported severe security staffing shortages (IBM Security, 2024), and organizations without dedicated security teams have fewer resources for evaluating AI-specific controls. The modular reference implementation and commodity-hardware prompt benchmark lower the barrier to evaluating one part of this defense stack. Production use of the full architecture still requires engineering, domain calibration, and—in the attestation path—TEE-capable infrastructure.

6. Conclusions

This work presents **EmbedGuard**, a RAG-defense reference architecture that proposes four-layer anomaly-signal fusion combined with a hardware-rooted embedding-provenance protocol; the released package implements the fusion logic and a software simulation of the protocol. The literature scan in Section 2.4 found close work on layered prompt-injection defense, defense orchestration, and software-signed embedding provenance; the contribution

claimed here is the specific integration of joint four-layer scoring with a TEE trust root, not the invention of layered defense or provenance in isolation. The repository releases the 83-pattern prompt detector, its 135-sample benchmark, and the correlation implementation so that the reproducible subset can be extended, attacked, and improved.

As retrieval-augmented generation systems span multiple attack surfaces, security evaluation should preserve evidence across the prompt, provenance, retrieval, and output stages rather than assume one detector is sufficient. EmbedGuard makes that integration testable, but the open artifact does not establish superiority over current defenses or production-grade full-pipeline latency.

For vectors injected or modified outside the attested ingestion path, the target attestation scheme replaces anomaly inference with verification of a provenance claim, subject to the TEE assumptions and CVEs in Section 4.5. For poisoned documents admitted through that path it provides no signal, which the open PoisonedRAG evaluation confirms (0/50; Section 4.3). The published Tier-1 evaluation reports 94.7% detection of optimization-based attacks and 89.3% under adaptive attack, with an 18.4 percentage point cross-layer ablation gain. Separately, the current open Tier-2 benchmark observes 30/30 prompt attacks detected and 0/105 benign queries flagged, with two-sided 95% Wilson lower bounds of 88.6% and 96.5%, respectively. On the open PoisonedRAG evaluation, a clean-calibrated distribution-distance detector flagged 50/50 attacks with five planted passages per target and 7/50 with one, at 6/200 flagged clean queries. The repository does not reproduce the Tier-1 deployment, evaluate an adaptive corpus-poisoning attacker, or evaluate established jailbreak suites.

The operational modes describe how organizations could adapt the architecture to different risk tolerances—particularly in healthcare, financial services, and legal settings where integrity failures affect safety, compliance, or professional liability. Beyond sector-specific applications, EmbedGuard exposes a testable reference architecture and a low-cost prompt-layer benchmark. That is a starting point for evaluation, not evidence that the full framework is production-ready or state of the art across deployment contexts.

Future Work: The most direct next experiments are an adaptive attacker that knows the distance detector, full-corpus retrieval, per-passage rather than retrieved-set statistics to address single-passage poisoning, and an evaluation of all four layers with the fused decision on the same workload. Beyond these, we will extend the framework to address: (1) multi-modal RAG systems with image/audio retrieval; (2) federated retrieval architectures with distributed trust; (3) continuous learning scenarios with evolving knowledge bases; and (4) availability attacks through retrieval denial pattern detection.

Acknowledgments

The author acknowledges the anonymous reviewers for their constructive feedback on earlier versions of this manuscript.

Data Availability

The reference implementation and Tier-2 evaluation materials are available with FAIR-oriented archival:

Primary Repository: <https://github.com/neerazz/embedguard> (MIT License)

Archived Versions: Zenodo concept DOI [10.5281/zenodo.18364919](https://doi.org/10.5281/zenodo.18364919) (resolves to latest archive); v1.0.0 = [10.5281/zenodo.18364920](https://doi.org/10.5281/zenodo.18364920) (version of record), v1.1.0 = [10.5281/zenodo.21280092](https://doi.org/10.5281/zenodo.21280092) (post-publication maintenance). Repository v1.3.0 / manuscript v3.2 adds the Tier-2b evaluation and is archived as v1.3.0 = [10.5281/zenodo.23032135](https://doi.org/10.5281/zenodo.23032135).

Contents:

- **Source Code:** Complete EmbedGuard framework implementation (Python 3.10+)
- **Benchmark Inputs:** locally curated Natural Questions-style (N=50), HotpotQA-style (N=25), and MS-MARCO-style (N=25) files; upstream subset provenance is not established
- **Injection Regression Set:** 35 rows: 30 attacks spanning 25 attack categories plus 5 benign controls, all with ground-truth labels
- **Evaluation Scripts:** Benchmark runner plus count-based Wilson interval analysis
- **Detection Patterns:** 83 patterns in `embedguard/prompt_detector/__init__.py` (INJECTION_PATTERNS)
- **Tier-2b Scripts and Results:** `scripts/fetch_tier2b_data.py` (pinned upstream revisions, SHA-256 checked), `scripts/tier2b_poisonedrag.py`, `scripts/run_tier2b.sh`, and six result files with per-query rows under `results/`. The PoisonedRAG passages (MIT) and BEIR Natural Questions (CC BY-SA 4.0) are downloaded from their sources, not redistributed
- **Container Recipe:** A Dockerfile for building a local image; no pre-built image is claimed

Reproducibility Commands:

```
# Option 1: One command (venv + tests + benchmark report)
git clone https://github.com/neerazz/embedguard.git && cd embedguard && ./reproduce.sh

# Option 2: Docker (build locally from the included Dockerfile)
docker build -t embedguard . && docker run --rm embedguard python examples/run_benchmarks.py

# Tier 2b (needs the neural extra and about 765 MB of inputs)
pip install -e ".[neural]" pyarrow
python scripts/fetch_tier2b_data.py --data-dir data/tier2b
scripts/run_tier2b.sh data/tier2b
```

Determinism Note: Tier-2 classification uses fixed inputs and deterministic regex/normalization logic; no random seed is required for the reported classification counts. Tier 2b fixes seed 12 for the pool and splits; embedding values can differ in the last digits across hardware back ends, so the result files record the device and library versions. Latency measurements vary by host and load. Appendix A separates the published Tier-1 environment from the commodity Tier-2 path.

References

Addison P, Nguyen MH, Medan T, Shah J, Manzari MT, McElrone B, Lalwani L, More A, Sharma S, Roth HR, Yang I, Chen C, Xu D, Cheng Y, Feng A, Xu Z. 2024. C-FedRAG: A Confidential Federated Retrieval-Augmented Generation System. arXiv preprint arXiv:2412.13163. DOI: 10.48550/arXiv.2412.13163.

AMD. 2020. AMD SEV-SNP: Strengthening VM Isolation with Integrity Protection and More. AMD White Paper (January 2020). Available: <https://www.amd.com/content/dam/amd/en/documents/epyc-business-docs/white-papers/SEV-SNP-strengthening-vm-isolation-with-integrity-protection-and-more.pdf> (accessed 2026-01-24).

AMD. 2025a. SEV-SNP Firmware Vulnerabilities. AMD Security Bulletin AMD-SB-3007. Available: <https://www.amd.com/en/resources/product-security/bulletin/amd-sb-3007.html> (accessed 2026-01-24).

AMD. 2025b. Guest Memory Vulnerabilities. AMD Security Bulletin AMD-SB-3011. Available: <https://www.amd.com/en/resources/product-security/bulletin/amd-sb-3011.html> (accessed 2026-01-24).

Carlini N, Nasr M, Choquette-Choo CA, Jagielski M, Gao I, Awadalla A, Koh PW, Ippolito D, Lee K, Tramer F, Schmidt L. 2023. Are aligned neural networks adversarially aligned? In: Advances in Neural Information Processing Systems 36 (NeurIPS 2023):61478-61500. arXiv preprint arXiv:2306.15447. DOI: 10.52202/075280-2687.

Chao P, Debenedetti E, Robey A, Andriushchenko M, Croce F, Sehwag V, Dobriban E, Flammarion N, Pappas GJ, Tramer F, Hassani H, Wong E. 2024. JailbreakBench: An Open Robustness Benchmark for Jailbreaking Large Language Models. In: Advances in Neural Information Processing Systems 37 (NeurIPS 2024, Datasets and Benchmarks Track). arXiv preprint arXiv:2404.01318.

Chaudhari H, Severi G, Abascal J, Suri A, Jagielski M, Choquette-Choo CA, Nasr M, Nita-Rotaru C, Oprea A. 2024. Phantom: General Backdoor Attacks on Retrieval Augmented Language Generation. arXiv preprint arXiv:2405.20485. DOI: 10.48550/arXiv.2405.20485.

Cheng Z, Sun J, Gao A, Quan Y, Liu Z, Hu X, Fang M. 2025. Secure Retrieval-Augmented Generation against Poisoning Attacks. In: Proceedings of IEEE BigData 2025. arXiv preprint arXiv:2510.25025. DOI: 10.48550/arXiv.2510.25025.

Chrapek M, Vahldiek-Oberwagner A, Spoczynski M, Constable S, Vij M, Hoefler T. 2024. Fortify Your Foundations: Practical Privacy and Security for Foundation Model Deployments in the Cloud. arXiv preprint arXiv:2410.05930. DOI: 10.48550/arXiv.2410.05930.

Geng R, Zou W, Wang B, Jia J. 2024. PoisonedRAG adversarial texts for Natural Questions (results/adv_targeted_results/nq.json). GitHub repository sleepeer/PoisonedRAG, commit f660d72, MIT License. Available: <https://github.com/sleepeer/PoisonedRAG> (accessed 2026-09-28).

IBM Security, Ponemon Institute. 2024. Cost of a Data Breach Report 2024. IBM Corporation. Available: <https://www.ibm.com/reports/data-breach> (accessed 2026-01-24; the URL now serves later editions).

Kim M, Lee H, Koo H. 2025a. Rescuing the Unpoisoned: Efficient Defense against Knowledge Corruption Attacks on RAG Systems. In: Proceedings of the Annual Computer Security Applications Conference (ACSAC 2025). arXiv preprint arXiv:2511.01268. DOI: 10.48550/arXiv.2511.01268.

Kim S, Kim J, Jeon Y, Lee GG. 2025b. Safeguarding RAG Pipelines with GMTP: A Gradient-based Masked Token Probability Method for Poisoned Document Detection. In: Findings of the Association for Computational Linguistics: ACL 2025, pp. 24597-24614. DOI: 10.18653/v1/2025.findings-acl.1263.

Kwiatkowski T, Palomaki J, Redfield O, Collins M, Parikh A, Alberti C, Epstein D, Polosukhin I, Devlin J, Lee K, Toutanova K, Jones L, Kelcey M, Chang MW, Dai AM, Uszkoreit J, Le Q, Petrov S. 2019. Natural Questions: A Benchmark for Question Answering Research. Transactions of the Association for Computational Linguistics 7:453-466. DOI: 10.1162/tacl_a_00276.

Lee D, Lee J, Ha JW, Kim JH, Lee SW, Lee H, Song HO. 2023. Query-Efficient Black-Box Red Teaming via Bayesian Optimization. In: Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL 2023). arXiv preprint arXiv:2305.17444. DOI: 10.48550/arXiv.2305.17444.

Lewis P, Perez E, Piktus A, Petroni F, Karpukhin V, Goyal N, Kuttler H, Lewis M, Yih W, Rocktaschel T, Riedel S, Kiela D. 2020. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. In: Advances in Neural Information Processing Systems 33 (NeurIPS 2020):9459-9474. DOI: 10.48550/arXiv.2005.11401.

Liu Y, Deng G, Li Y, Wang K, Wang Z, Wang X, Zhang T, Liu Y, Wang H, Zheng Y, Zhang LY, Liu Y. 2023. Prompt Injection attack against LLM-integrated Applications. arXiv preprint arXiv:2306.05499. DOI: 10.48550/arXiv.2306.05499.

Menlo Ventures. 2024. 2024: The State of Generative AI in the Enterprise. Available: <https://menlovc.com/2024-the-state-of-generative-ai-in-the-enterprise/> (accessed 2026-07-09).

Misono M, Stavrakakis D, Santos N, Bhatotia P. 2024. Confidential VMs Explained: An Empirical Analysis of AMD SEV-SNP and Intel TDX. *Proceedings of the ACM on Measurement and Analysis of Computing Systems* 8(3):1-42. DOI: 10.1145/3700418.

MITRE. 2025. MITRE ATLAS: Adversarial Threat Landscape for Artificial-Intelligence Systems — Techniques AML.T0064, AML.T0066, AML.T0070, AML.T0071. Available: <https://atlas.mitre.org/> (accessed 2026-07-09).

NVD. 2025. CVE-2025-32711: M365 Copilot Information Disclosure Vulnerability (EchoLeak). National Vulnerability Database. Available: <https://nvd.nist.gov/vuln/detail/CVE-2025-32711> (accessed 2026-07-09).

OWASP. 2025. OWASP Top 10 for Large Language Model Applications 2025 — LLM01: Prompt Injection; LLM04: Data and Model Poisoning; LLM08: Vector and Embedding Weaknesses. Available: <https://genai.owasp.org/> (accessed 2026-07-09).

Pallerla P, Bhukya WN, Vemula B, Kodi CR. 2026. Adaptive Defense Orchestration for RAG: A Sentinel-Strategist Architecture against Multi-Vector Attacks. *arXiv preprint arXiv:2604.20932*.

Ramakrishnan B, Balaji A. 2025. Securing AI Agents Against Prompt Injection Attacks. *arXiv preprint arXiv:2511.15759*.

Rehberger J. 2024. Spyware Injection Into Your ChatGPT's Long-Term Memory (SpAIware). Embrace The Red. Available: <https://embracethered.com/blog/posts/2024/chatgpt-macos-app-persistent-data-exfiltration/> (accessed 2026-07-09).

Rehberger J. 2025. Hacking Gemini's Memory with Prompt Injection and Delayed Tool Invocation. Embrace The Red. Available: <https://embracethered.com/blog/posts/2025/gemini-memory-persistence-prompt-injection/> (accessed 2026-07-09).

Reimers N, Gurevych I. 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In: *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP-IJCNLP 2019)*:3980-3990. DOI: 10.18653/v1/D19-1410.

RoyChowdhury A, Luo M, Sahu P, Banerjee S, Tiwari M. 2024. ConfusedPilot: Confused Deputy Risks in RAG-based LLMs. *arXiv preprint arXiv:2408.04870*. DOI: 10.48550/arXiv.2408.04870.

Saleem G, Ahmed N, Zaman MI, Hassan A, Mujahid U. 2026. A Layered Security Framework Against Prompt Injection in RAG-Based Chatbots. *arXiv preprint arXiv:2606.19660*.

Shen Z, Imana B, Wu T, Xiang C, Mittal P, Korolova A. 2025. ReliabilityRAG: Effective and Provably Robust Defense for RAG-based Web-Search. In: *Advances in Neural Information Processing Systems 38 (NeurIPS 2025)*. *arXiv preprint arXiv:2509.23519*. DOI: 10.48550/arXiv.2509.23519.

Sun X, Li X, Meng Y, Ao X, Lyu L, Li J, Zhang T. 2023. Defending against Backdoor Attacks in Natural Language Generation. In: *Proceedings of the AAAI Conference on Artificial Intelligence* 37(4):5257-5265. DOI: 10.1609/aaai.v37i4.25656.

Tan X, Luan H, Luo M, Sun X, Chen P, Dai J. 2025. RevPRAG: Revealing Poisoning Attacks in Retrieval-Augmented Generation through LLM Activation Analysis. In: *Findings of the Association for Computational Linguistics: EMNLP 2025*. *arXiv preprint arXiv:2411.18948*. DOI: 10.48550/arXiv.2411.18948.

Thakur N, Reimers N, Rücklé A, Srivastava A, Gurevych I. 2021. BEIR: A Heterogenous Benchmark for Zero-shot Evaluation of Information Retrieval Models. In: *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks 1 (NeurIPS 2021)*. *arXiv preprint arXiv:2104.08663*. DOI: 10.48550/arXiv.2104.08663.

Vassilev A, Oprea A, Fordyce A, Anderson H. 2025. Adversarial Machine Learning: A Taxonomy and Terminology of Attacks and Mitigations. NIST AI 100-2e2025. DOI: 10.6028/NIST.AI.100-2e2025.

Wanger J. 2026. VectorSmuggle: Steganographic Exfiltration in Embedding Stores and a Cryptographic Provenance Defense. arXiv preprint arXiv:2605.13764.

Wilson EB. 1927. Probable Inference, the Law of Succession, and Statistical Inference. Journal of the American Statistical Association 22(158):209-212. DOI: 10.1080/01621459.1927.10502953.

Xiang C, Wu T, Zhong Z, Wagner D, Chen D, Mittal P. 2024. Certifiably Robust RAG against Retrieval Corruption. arXiv preprint arXiv:2405.15556. DOI: 10.48550/arXiv.2405.15556.

Yao H, Shi H, Chen Y, Jiang Y, Wang C, Qin Z. 2025. ControlNET: A Firewall for RAG-based LLM System. arXiv preprint arXiv:2504.09593. DOI: 10.48550/arXiv.2504.09593.

Zhang B, Xin H, Fang M, Liu Z, Yi B, Li T, Liu Z. 2025. Traceback of Poisoning Attacks to Retrieval-Augmented Generation. In: Proceedings of the ACM Web Conference 2025 (WWW 2025), pp. 2085-2097. DOI: 10.1145/3696410.3714756.

Zhou H, Lee KH, Zhan Z, Chen Y, Li Z, Wang Z, Haddadi H, Yilmaz E. 2025a. TrustRAG: Enhancing Robustness and Trustworthiness in Retrieval-Augmented Generation. arXiv preprint arXiv:2501.00879. DOI: 10.48550/arXiv.2501.00879.

Zhou P, Feng Y, Yang Z. 2025b. Provably Secure Retrieval-Augmented Generation. arXiv preprint arXiv:2508.01084. DOI: 10.48550/arXiv.2508.01084.

Zou A, Wang Z, Carlini N, Nasr M, Kolter JZ, Fredrikson M. 2023. Universal and Transferable Adversarial Attacks on Aligned Language Models. arXiv preprint arXiv:2307.15043. DOI: 10.48550/arXiv.2307.15043.

Zou W, Geng R, Wang B, Jia J. 2025. PoisonedRAG: Knowledge Corruption Attacks to Retrieval-Augmented Generation of Large Language Models. In: Proceedings of the 34th USENIX Security Symposium. arXiv preprint arXiv:2402.07867. DOI: 10.48550/arXiv.2402.07867.

Appendix A: Experimental Infrastructure and Reproducibility

A.1 Tier-1 Hardware Configuration (Version-of-Record Reference)

The configuration below is reported for the published Tier-1 evaluation. It is not recreated by reproduce.sh and should not be read as the hardware used for the current open benchmark.

Component	Specification
Processor	AMD EPYC 7542, 32 cores, 2.9GHz base / 3.4GHz boost
Memory	256GB DDR4-3200 ECC Registered
TEE Platform	AMD SEV-SNP
TEE Firmware	AMD-SP firmware version 1.55.x
Storage	Samsung PM9A3 NVMe SSD, 1.92TB
Network	Mellanox ConnectX-6 100GbE

A.2 Tier-1 Software Stack (Version-of-Record Reference)

Component	Version
Operating System	Ubuntu 22.04.3 LTS
Linux Kernel	6.5.0-generic (SEV-SNP enabled)
Python	3.10.12
PyTorch	2.1.0+cu118
Transformers	4.35.2
Sentence-Transformers	2.2.2
Embedding Model	sentence-transformers/all-mpnet-base-v2
FAISS	1.7.4 (GPU)
NumPy	1.24.3
scikit-learn	1.3.2

A.3 Tier-1 TEE Configuration (Version-of-Record Reference)

Parameter	Setting
SEV-SNP Policy	0x30000 (debug disabled, migration disabled)
Attestation	AMD Key Distribution Service (KDS)
Guest VMPL	Level 0 (highest privilege within guest)
Memory Encryption	AES-128-XTS

A.4 Reproducibility Checklist

- ☒ Code available at: <https://github.com/neerazz/embedguard>
- ☒ Archived versions with DOIs: Zenodo concept 10.5281/zenodo.18364919 (v1.0.0: 18364920, v1.1.0: 21280092, v1.3.0: 23032135)
- ☒ Benchmark datasets included (data/ directory)
- ☒ Tier-1 hardware/software specifications documented as version-of-record reference
- ☒ Tier-2 dependency ranges declared (pyproject.toml, requirements.txt)
- ☒ Pattern-based classifier (83 patterns) included in source
- ☒ Benchmark runner script (run_benchmarks.py) included
- ☒ Count-based uncertainty analysis (scripts/statistical_tests.py) included
- ☐ Tier-1 production corpus and TEE execution environment publicly reproducible
- ☒ Published corpus-poisoning attack (PoisonedRAG, NQ) evaluated on the provenance and retrieval layers with a held-out split (Section 4.3)
- ☐ Adaptive attacker against the distribution-distance detector evaluated
- ☐ All four layers and the fused decision evaluated on one corpus-poisoning workload
- ☐ JailbreakBench or equivalent external adversarial suite evaluated

A.5 Evaluation Datasets

Dataset	Size	Domain	Source
Natural Questions-style	50 samples	Curated open-domain QA	Local file labeled NQ; no upstream IDs
HotpotQA-style	25 samples	Curated multi-hop reasoning	Local file labeled HotpotQA; no extraction manifest
MS-MARCO-style	25 samples	Curated web-search QA	Local file labeled MS-MARCO; no extraction manifest

Dataset	Size	Domain	Source
Injection Attacks	30 attacks	25 attack categories	Curated adversarial dataset
Benign Control	5 samples	Mixed domains	Baseline verification set
Total	135 samples	Multi-domain	Reproducible benchmark

Attack Categories (25 types): direct_instruction, jailbreak_attempt, instruction_smuggling, context_manipulation, prompt_leaking, role_manipulation, indirect_injection, unicode_obfuscation, base64_encoding, delimiter_confusion, xml_injection, markdown_injection, developer_mode, hypothetical_framing, fictional_framing, translation_attack, repetition_attack, authority_claim, emotional_manipulation, rag_specific (3), composite_attack (2), subtle_manipulation (2), multi_turn_setup, payload_splitting, virtualization.

A.6 Compute Requirements

Experiment	Execution	Recorded detector time
Pattern Matching (135 samples)	Single-threaded	~11.4ms summed detector time on recorded host
Per-Query Latency	Standard CPU	0.085ms aggregate mean on recorded host

Note: The current benchmark focuses on the prompt classifier layer and requires no GPU. Optional neural detection, embedding generation, and vector search may use accelerator-specific extras in larger RAG deployments; the released pattern-only benchmark runs on a standard CPU.